

Operate Kubernetes workloads

extend the platform with the operator pattern!

Paolo Patierno

Senior Principal Software Engineer @Red Hat



CLOUD NATIVE
COMPUTING NAPOLI

About me

apiVersion: v1

kind: SeniorPrincipalSoftwareEngineer

metadata:

name: Paolo Patierno

namespace: Red Hat, Messaging & Data Streaming

labels:

cncf/maintainer: Strimzi

cncf/maintainer: true

eclipse/committer: Vert.x, Hono & Paho

microsoft/mvp: Azure

annotations:

family: dad of two, husband of one

sports: running, swimming, motogp, vr46, formula1, ferrari, ssc napoli

community: cncf napoli, devday

spec:

replicas: 1

containers:

- **image:** patiernohub.io/paolo:latest



@ppatierno



CLOUD NATIVE
COMPUTING NAPOLI

Kubernetes workloads ...



How does Kubernetes handle scaling, rollout, batch execution and so on?



Not just pods ...

- Don't use Pod(s) ... let's use something more sophisticated!
- ReplicaSet
 - Guarantees a specific number of running replicas
 - Spins pods, based on a template, if there are not enough ...
 - ... or deletes pods if too many
- Deployment
 - It's based on ReplicaSet (used to run replicas)
 - Adds extra layer for rollout and rollback
- ... and more with StatefulSet, Job, DaemonSet ...

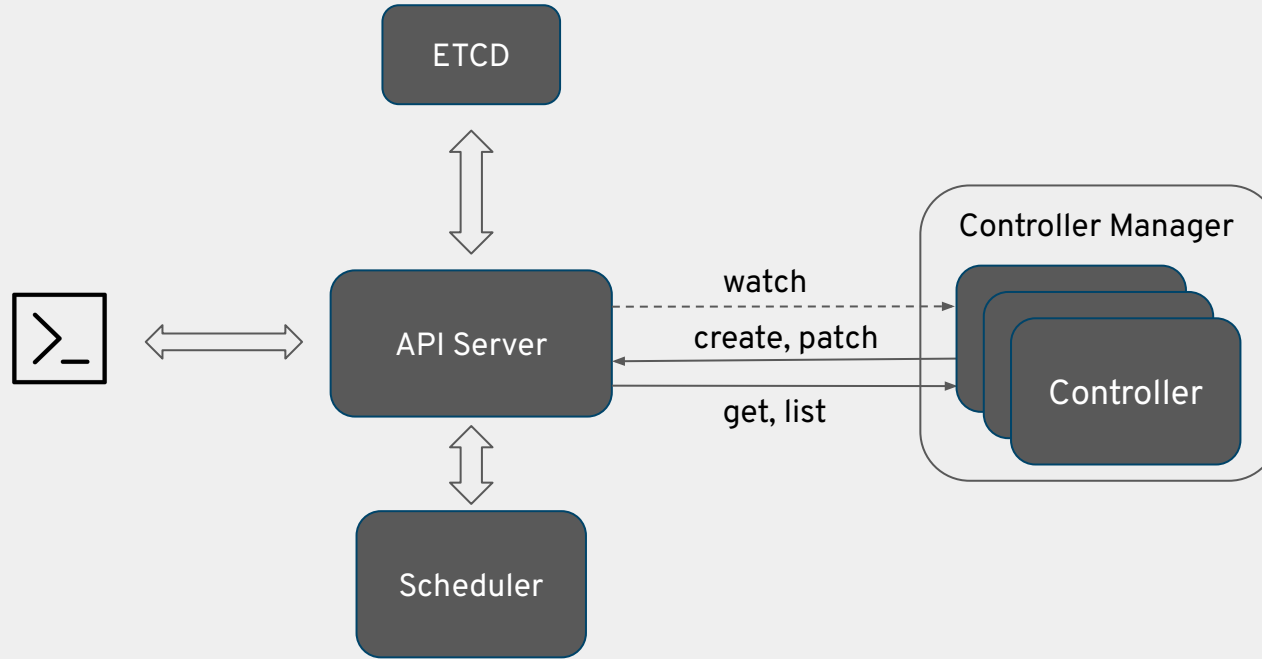




How does it work?
Let's use a controller!!!
.... But not this one ;-)

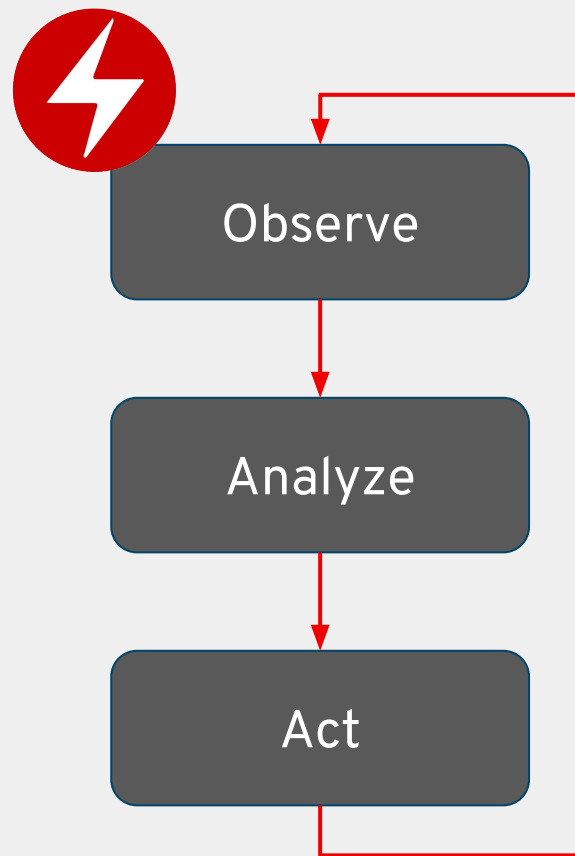


Kubernetes Control Plane



Reconcile Loop

- **Observe**
 - Watch for resource/object creation or changes
- **Analyze**
 - Check that the resource/object desired state (“spec”) reflects the current state on the cluster
- **Act**
 - Makes the needed changes



ReplicaSet Controller

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
```

my-replicaset-bf5zv

my-replicaset-1tf5a


my-replicaset-gb65f



ReplicaSet Controller

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
```

replicas: 5



my-replicaset-bf5zv

my-replicaset-1tf5a

my-replicaset-gb65f

my-replicaset-5tfgs

my-replicaset-rf43g



ReplicaSet Controller

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  containers:
    apiVersion: v1
    kind: Pod
    metadata:
      name: my-pod-2
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
          image: quay.
```

my-replicaset-bf5zv

my-replicaset-1tf5a

my-replicaset-gb65f

~~my-pod-1~~

~~my-pod-2~~



ReplicaSet Controller

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  containers:
    apiVersion: v1
    kind: Pod
    metadata:
      name: my-pod-2
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
          image: quay.
```

my-pod-1

my-pod-2



ReplicaSet Controller


```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
```

my-pod-1

my-pod-2

my-replicaset-65rt3






How to automate
operating complex
applications?
The Operator pattern!

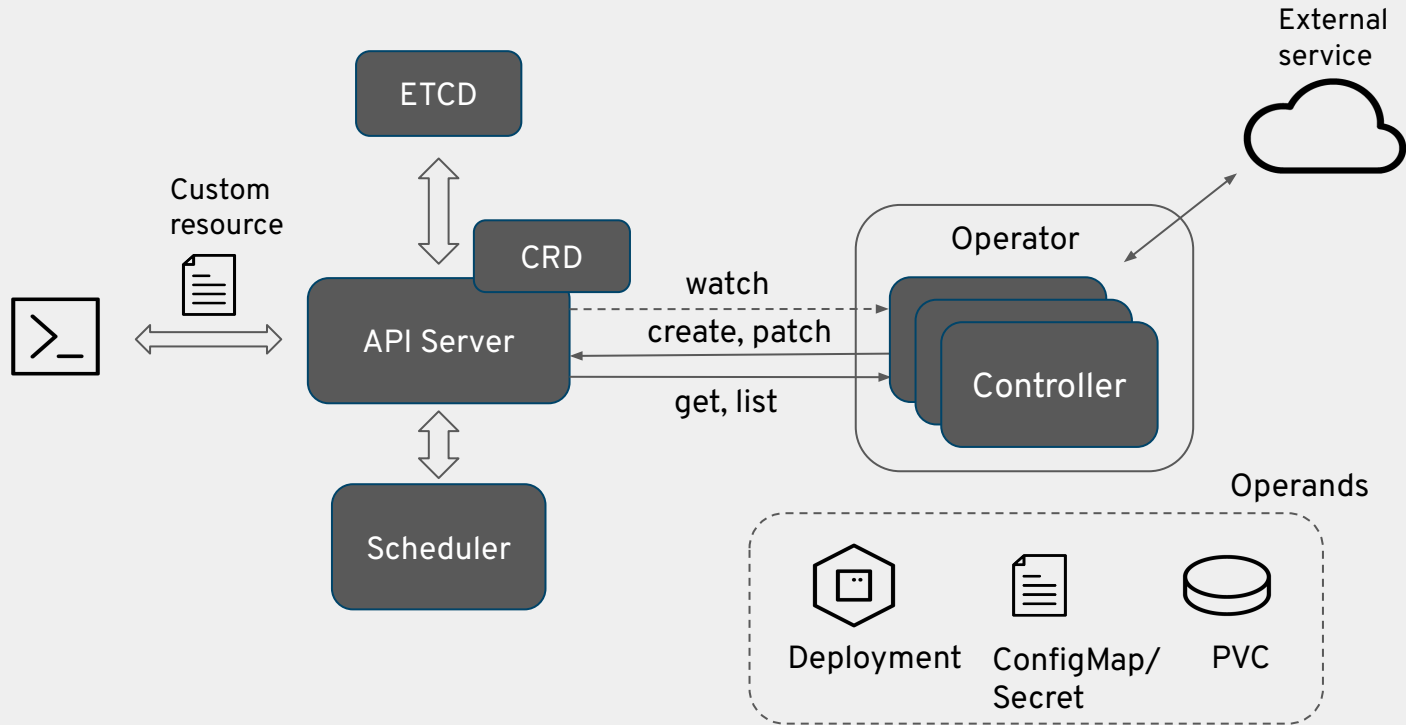


Operator

- It's yet another containerized application!
- Has the knowledge of a specific business domain
- Manage the application lifecycle
- Leverage CRDs (Custom Resource Definition) to extend API server
- Takes care of one (or more) custom resources/objects
 - By having a controller for each resource
 - Creating native Kubernetes resources ... aka “operands”
 - Leveraging built-in controllers via API server interaction



Operator Pattern



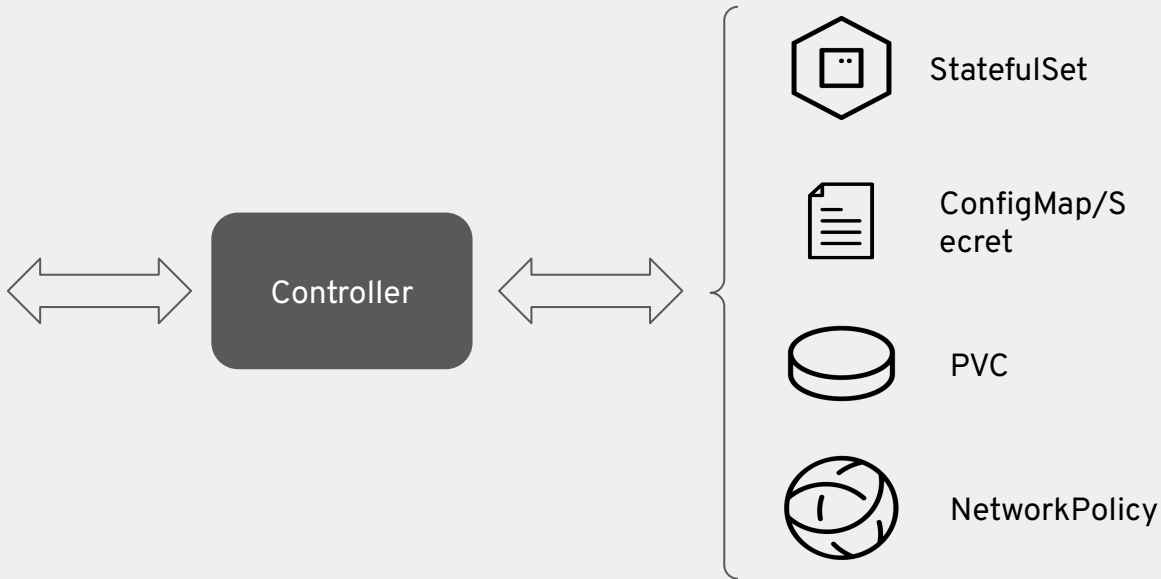
From the Custom Resource Definition ...

- Apache Kafka is a ... Kubernetes resource!
- Declare a new Kubernetes “kind”
 - Group
 - Versions
- Define the new “kind” structure using an OpenAPI schema
 - Spec
 - status

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: kafkas.kafka.strimzi.io
spec:
  group: kafka.strimzi.io
  names:
    kind: Kafka
    listKind: KafkaList
  ...
  versions:
  - name: v1beta2
    schema:
      openAPIV3Schema:
        type: object
        properties:
          spec:
            # spec definition with for
            # the custom resource
            ...
          status:
            # status definition reported back
            # in the custom resource
```

... to the Custom Resource

```
apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    version: 3.1.0
    replicas: 3
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
      - name: tls
        port: 9093
        type: internal
        tls: true
    config:
      default.replication.factor: 3
      min.insync.replicas: 2
      inter.broker.protocol.version: "3.1"
      ...
    storage:
      type: ephemeral
  zookeeper:
    replicas: 3
    storage:
      type: ephemeral
  status:
    ...
    ...
```





VS



Why? What
about Helm
Charts?



CLOUD NATIVE
COMPUTING NAPOLI

Helm Charts

- “Package manager” for Kubernetes
 - Charts = template + values
- Rely on Kubernetes built-in resources (i.e. Deployment, ConfigMap, ...)
- Simplify to write YAMLS with parameters via templating
- Simplify day-1 operation, for deploying applications
- Key problems fixed
 - Deploy same application with different configuration
 - Deploy same application on different environments



Operators

- Control “life cycle” of Kubernetes workloads
 - Operator = CRD(s) + Controller(s)
- Extend the Kubernetes API with CRDs (Custom Resource Definitions)
- Simplify to write one (or a few) “custom resource” related YAMLs
- Acting since day-1 to day-2 operation from deployment to upgrades, through manage
- ... deployable via an Helm Charts :-)
 - Helm guarantees CRDs are installed before operator
 - Operator configurable via values or ConfigMap



Thank you!



CLOUD NATIVE
COMPUTING NAPOLI